

CS3245

# Information Retrieval

# 8

Lecture 8: A complete search system –  
Scoring and results assembly



Live Q&A  
<https://pollev.com/jin>

# Last Time: tf-idf weighting



- The tf-idf weight of a term is the product of its *tf* weight and its *idf* weight.

$$w_{t,d} = (1 + \log \text{tf}_{t,d}) \times \log(N / \text{df}_t)$$

- Best known weighting scheme in information retrieval
  - One of the easy but important things you should remember for IR
  - Increases with the number of occurrence within a document
  - Increases with the rarity of the term in the collection

# Last Time: Vector Space Model

- Key idea 1: represent both  $d$  and  $q$  as vectors
- Key idea 2: Rank documents according to their proximity (similarity) to the query in this space

$$\cos(\vec{q}, \vec{d}) = \frac{\vec{q} \bullet \vec{d}}{|\vec{q}| |\vec{d}|} = \frac{\sum_{i=1}^{|V|} q_i d_i}{\sqrt{\sum_{i=1}^{|V|} q_i^2} \sqrt{\sum_{i=1}^{|V|} d_i^2}}$$

$\cos(q, d)$  is the cosine similarity of  $q$  and  $d$  ... or, equivalently, the cosine of the angle between  $q$  and  $d$ .

# Computing cosine scores, redux

COSINESCORE( $q$ )

$$\cos(\vec{q}, \vec{d}) = \frac{\sum_{i=1}^{|\mathcal{V}|} q_i d_i}{\sqrt{\sum_{i=1}^{|\mathcal{V}|} q_i^2} \sqrt{\sum_{i=1}^{|\mathcal{V}|} d_i^2}}$$

*Consider only the terms  
appearing in both  $q$  and  $d$ .*

```

1  float Scores[N] = 0
2  float Length[N]
3  for each query term  $t$ 
4  do calculate  $w_{t,q}$  and fetch postings list for  $t$ 
5      for each pair( $d, tf_{t,d}$ ) in postings list
6          do  $Scores[d] + = w_{t,d} \times w_{t,q}$ 
7  Read the array  $Length$ 
8  for each  $d$ 
9      do  $Scores[d] = Scores[d] / Length[d]$ 
10 return Top  $K$  components of  $Scores[]$ 
```

Dot product

*Normalize by the (pre-computed)  
document length only.*

Normalization

# Today



## Goal


- Speeding up and shortcutting ranking
- Incorporating additional ranking information into VSM

# Efficient cosine ranking



- Key observations
  - Users **only checks the top results.**
  - There are probably too many (relevant) documents in the first place.
- Given a collection of  $N$  documents and a query
  - Find  $K$  ( $\ll N$ ) docs that are **(likely to be)** the "nearest" to the query based on cosine similarity.
- Efficient ranking
  - Simplify the processing
  - Possibly less accurate / exact

# Faster cosine: unweighted query

- To simplify the computation of a single cosine, we can...
- Assume each query term has weight 1
  - i.e.,  $w_{t,q} = 1$  (no *tf*, nor *idf* factor; just Boolean presence)
  - Before:  $\text{Scores}[d] += w_{t,d} \times w_{t,q}$
  - After:  $\text{Scores}[d] += w_{t,d}$   *No expensive multiplication, only addition*
- But the bigger bottleneck is to process all  $N$  documents in the collection...

# Let's shrink the collection...



- Full collection =  $N$  documents
- Documents that do not contain any query terms have zero cosine values
  - Q: emperor
  - Doc1: queen, Doc2: the emperor, ...
  - $\text{Score}(Q, \text{Doc1}) = 0$
- Such documents can be safely ignored...Let's call the remaining collection of documents  $J$ .




# Optimizing the selection process

- What we need: Select  $K$  best out of  $J$ 
  - Typically,  $K \ll J$
  - Query: emperor
  - $J$  (i.e., docs containing emperor) = 1M, but  $K$  could be just 100
- Sort and output top  $K = O(J \log J + K)$
- Can we do better?

# Use heaps for selecting top $K$

- Heap = Binary tree in which  
each node's value  $>$  the values of its children
- Takes  $O(J)$  operations to construct, then each of  $K$  "winners" read off in  $O(\log J)$  steps =  $O(J + K * \log J)$
- For  $J = 1\text{M}$ ,  $K = 100$ , this is about 5% of the cost of sorting and outputting (with log base 2)



Blanks on slides, you may want to fill in

# Bottlenecks




- Primary computational bottleneck in scoring: cosine computation
- Can we avoid doing this computation for all docs in J?
  - Yes, we need to do some pruning.
- We may get it wrong sometimes but it is ok if we are not missing too many.
  - It is unlikely that the user really want **all** relevant documents.

# Generic approach



- Find a set  $A$  of contenders, with  $K < |A| \ll |J| \ll N$ 
  - $A$  does not necessarily contain the top  $K$ , but has many docs from among the top  $K$
  - Return the top  $K$  docs in  $A$
- Think of  $A$  as pruning non-contenders
- The same approach can also be used for other (non-cosine) scoring functions.



Blanks on slides, you may want to fill in



# Heuristic 1: Index elimination

- Basic algorithm: FastCosineScore of Fig 7.1 considers docs containing at least one query term (i.e., set J)
  - 4 for each query term  $t$
  - 5 do calculate  $w_{t,q}$  and fetch postings list for  $t$
  - 6 for each pair  $(d, tf_{t,d})$  in postings list
- J will be large and the computation will be slow if
- We can in fact ignore part of the index (i.e., postings lists) based on the query.

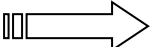
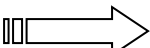
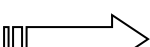
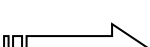
## 1a. High-idf query terms only

- E.g., given a query such as ***catcher in the rye*** only accumulate scores from *catcher* and *rye*
- It is usually not important to match **in** and **the** anyway since they have low idfs.
- **Benefit:**
  - Postings of low *idf* terms have many docs → these (many) docs get eliminated from set *A* of contenders
  - Similar in spirit to stop word removal

## 1b. Docs containing many query terms

- Any doc with at least one query term is a candidate from the top  $K$  output list, but ...
- For multi-term queries, only compute scores for docs containing several of the query terms
  - Say, at least 3 out of 4 query terms
- Easy to implement in postings traversal

# Example: Requiring 3 of 4 query terms

<b><i>Antony</i></b>		<table><tr><td>3</td><td>4</td><td>8</td><td>16</td><td>32</td><td>64</td><td>128</td><td></td></tr></table>	3	4	8	16	32	64	128	
3	4	8	16	32	64	128				
<b><i>Brutus</i></b>		<table><tr><td>2</td><td>4</td><td>8</td><td>16</td><td>32</td><td>64</td><td>128</td><td></td></tr></table>	2	4	8	16	32	64	128	
2	4	8	16	32	64	128				
<b><i>Caesar</i></b>		<table><tr><td>1</td><td>2</td><td>3</td><td>5</td><td>8</td><td>13</td><td>21</td><td>34</td></tr></table>	1	2	3	5	8	13	21	34
1	2	3	5	8	13	21	34			
<b><i>Calpurnia</i></b>		<table><tr><td>13</td><td>16</td><td>32</td><td></td><td></td><td></td><td></td><td></td></tr></table>	13	16	32					
13	16	32								


Scores only computed for docs 8, 16 and 32.



# Heuristic 2: Champion lists



- **Precompute** for each dictionary term  $t$ , the  $r$  docs of highest weight in  $t$ 's postings
  - Call this the champion list for  $t$   
(a.k.a. fancy list or top docs for  $t$ )



# Heuristic 2: Champion lists




- At query time, only compute scores for docs in the champion list of some query term
  - Pick the  $K$  top-scoring docs from amongst these
- Note that  $r$  has to be chosen at the indexing stage
  - Thus, it's possible that  $r < K$

# High and low lists



- For each term, we maintain two postings lists called *high* and *low*
  - Think of *high* as the champion



# High and low lists



- When traversing postings on a query, only traverse *high* lists first
  - If we get more than  $K$  docs, select the top  $K$  and stop
  - Else proceed to get docs from the *low* lists
- A means for segmenting index into two tiers

# Tiered indexes



- Generalizing high-low lists into tiers
- Break postings up into a hierarchy of lists

Most important


...

Least important

- Inverted index thus broken up into tiers of decreasing importance
- At query time, use only top tier unless insufficient to get K docs  
If so, drop to lower tiers




# Example tiered index



To think about:  
**What information  
would be useful to  
use to determine  
tiers?**

# Heuristic 3: Impact-ordered postings


- We only want to compute scores for docs for which  $wf_{t,d}$  is high enough
- We sort each postings list by  $wf_{t,d}$



## 3a. Early termination



**Antony**




8 (0.7)	16 (0.6)	3 (0.5)	4 (0.2)
---------	----------	---------	---------

- When traversing  $t$ 's postings (sorted by  $wf_{t,d}$ ), stop early after either
    - a fixed number of  $r$  docs
    - $wf_{t,d}$  drops below some threshold
- The score contribution ( $wf_{t,d} * wf_{t,q}$ ) is likely to be too low beyond these.
- Take the union of the resulting sets of docs
    - One set from the postings of each query term
  - Compute only the scores for docs in this union



## 3b. idf-ordered query terms

- Consider the postings of query terms in order of decreasing *idf*
  - Query: *story Caesar Antony*
  - Order of processing: *Antony Caesar story*
- Skip low-idf query terms completely (e.g., ignore *story*) ← Similar to 1a
- Move on to the next query term once the score contribution ( $wf_{t,d} * wf_{t,q}$ ) is low (e.g.,  $\leq 0.5$ )



E.g., if the query term weight of Anthony is **0.9**, skip to Caesar after checking the 3rd document.

# Heuristic 4:

## Cluster pruning – preprocessing



- Pick  $\sqrt{N}$  docs at random, call these *leaders*
- For other docs, pre-compute nearest leader
  - Docs attached to a leader are its followers
  - Likely: each leader has  $\sqrt{N}$  followers.


Why choose leaders at random?

- Fast
- Leaders reflect data distribution

# Cluster pruning visualization




1. Offline: Choose  $\sqrt{N}$  leaders



# Cluster pruning visualization



## 2. Associate documents to leaders to form clusters




# Cluster pruning – query processing

- Process a query as follows:
  - Given a query  $Q$ , find its nearest *leader*  $L$ .
  - Seek  $K$  nearest docs from among  $L$ 's followers (and  $L$  itself).

# Cluster pruning visualization



## 3. Online: Associate query to a leader (cluster)



# Clustering pruning variants



- Have each follower attached to  $b_1$  nearest leaders
- From query, find  $b_2$  nearest leaders and their followers
- $b_1$  affects preprocessing step at indexing time
- $b_2$  affects query processing step at run time

**To think about:** How do these parameters affect the retrieval results?

# Incorporating Additional Information: Static quality scores



- We want top-ranking documents to be both *relevant* and *authoritative*
  - *Relevance* is being modeled by cosine scores
  - *Quality* is typically a query-independent property of a document
- Examples of quality signals
  - Wikipedia among websites
  - Articles in certain newspapers
  - A paper with many citations
  - Many views, retweets, favs, bookmark saves
  - PageRank score

Quantitative



# Net score



- Assign to each document a quality score  $g(d)$  in  $[0,1]$ 
  - E.g., PageRank
- Combine cosine relevance and quality
$$\text{net-score}(q,d) = g(d) + \cos(q, d)$$
  - Can use some other linear combination than an equal weighting
- Now we seek the top  $K$  docs by net-score

# Incorporating Additional Information: Query term proximity



- Free text queries: just a set of terms typed into the query box – common on the web
- Users prefer docs where the query terms occur close to each other
- Let  $w$  be the smallest window in a document containing all query terms, e.g.,
  - Given the query *open day*:
    - For the document *open the next day*, the size of  $w$  is 4.
    - For the document *national day open house*, the size of  $w$  is 2.

# Query term proximity



- Collect candidates by running one or more queries to the indexes, and then rank.
- e.g., *NUS open day*
  1. Run it as a phrase query (e.g., using a positional index)
  2. If  $< K$  docs contain the phrase *NUS open day*, run the two phrase queries "*NUS open*" and "*open day*"
  3. If we still have  $< K$  docs, run the vector space query *NUS open day*
  4. Rank matching docs by vector space scoring combining all information (possibly including proximity score  $w$ )

# Incorporating Additional Information: Parametric and zone indexes



Documents often have multiple parts, with different semantics:

- Author, Title, Date of publication, etc.

These constitute the metadata about a document.

We sometimes wish to search by these metadata.

- E.g., find docs authored by T.S. Raffles in the year 1818, with *Dutch East India Company* in the title

# Fields



- **Year = 1818** is an example of a field
  - Also, **author = T.S. Raffles**, etc
  - with a **finite set** of possible values
- **Field or parametric index**
  - Postings for each field value
  - Sometimes build range (B-tree) trees (e.g., for dates)
- **Field query typically treated as conjunction**
  - find docs authored by **T.S. Raffles in the year 1818... =**
  - doc *must* be authored by T.S. Raffles AND in the year 1818.


# Zone



- A zone is a region of the doc that can contain an arbitrary amount of text e.g.,
  - Title
  - Abstract
  - References ...
- Build inverted indexes on zones as well to permit querying
  - E.g., find docs ... with Dutch East India Company in the title


# Two methods for zone indexing

## Alternative 1:

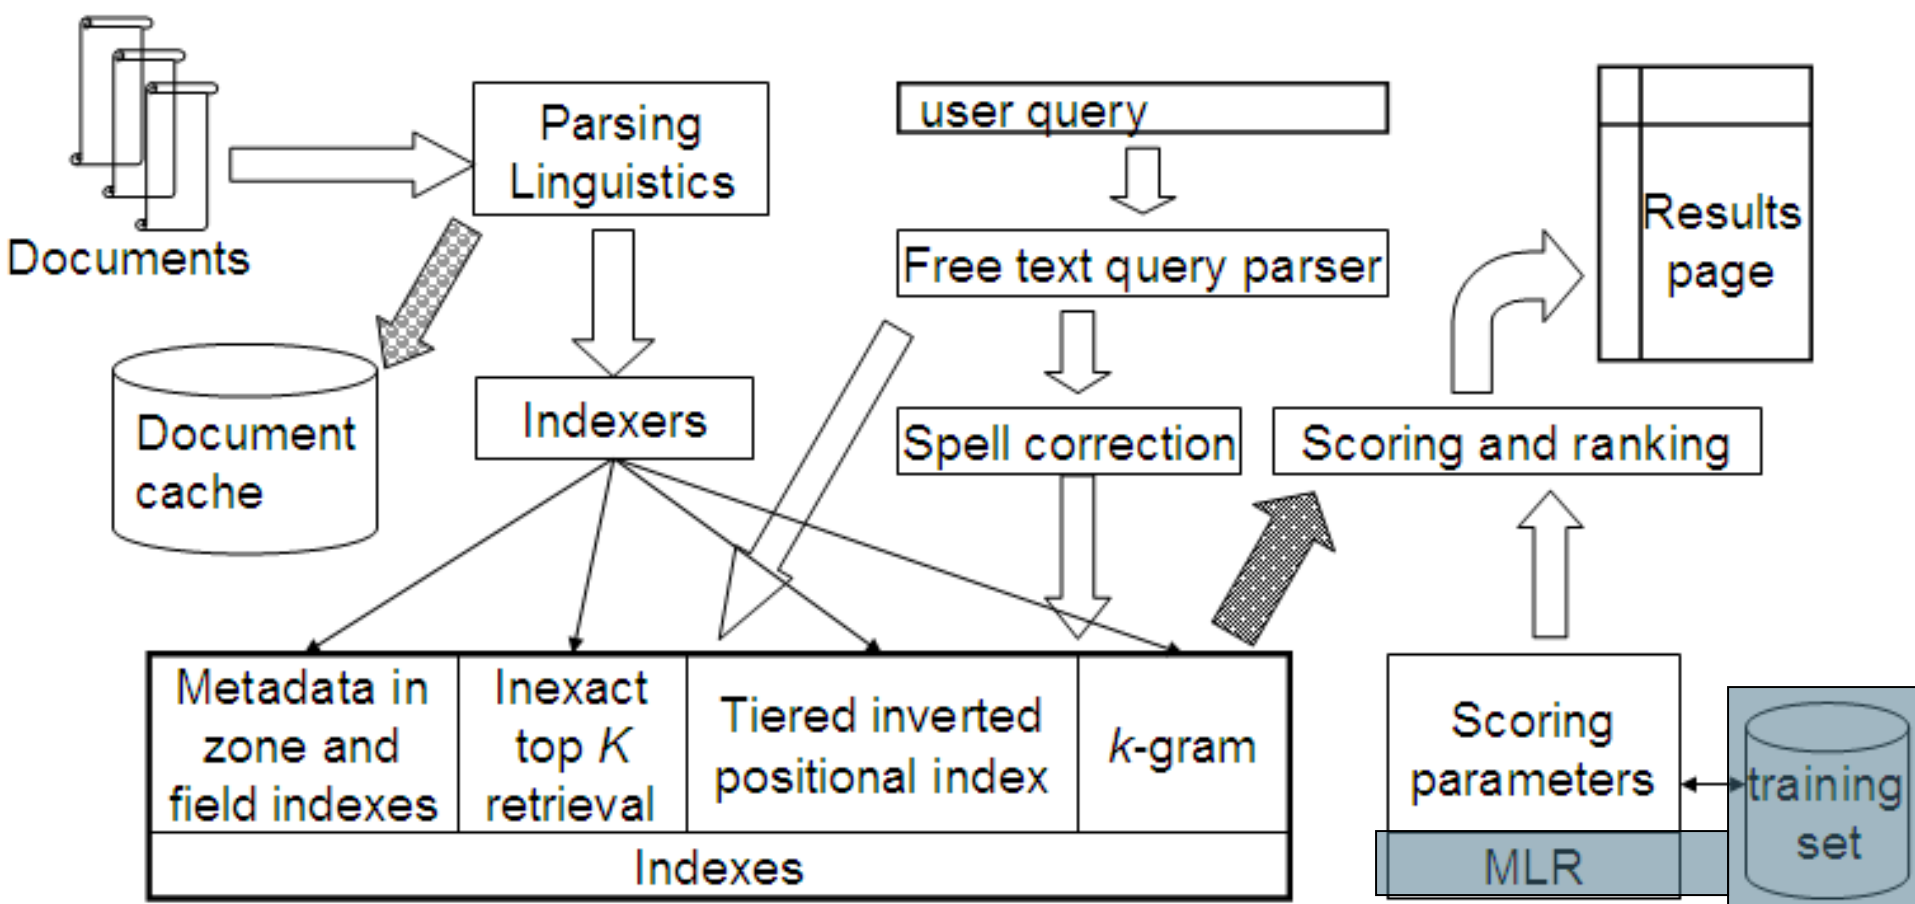


Encode zones in dictionary vs. postings.

## Alternative 2:



# Putting it all together



Won't be covering these blue modules in this course





# Summary

---

- Making the Vector Space Model more effective and efficient to compute
- Incorporating additional information

## Resources for today

- IIR 7, 6.1